

# Introduction to Networking and Systems Measurements

## Advanced Measurements



**Andrew W. Moore**  
andrew.moore@cl.cam.ac.uk


1

## Lessons from Lab1

- Ping isn't a perfect tool for latency measurements
- iperf isn't a perfect tool for bandwidth measurements
- Control, variability, accuracy, ....

2

## Example: Detecting Network Congestion




3

## How to control generated traffic?

- What is the packet format? (e.g. protocol, payload)
- How many packets?
- What is the packet size(s)?
- What is the average data rate?
- What is the peak data rate? (e.g. burst control)
- ...

4

## Traffic Generation Tools



\$\$\$\$\$, Hardware, high quality  
(Ixia, Spirent,...)

\$\$ Software/hardware based, medium quality  
(OSNT, MoonGen,...)

Commodity, Software, low quality  
(TCPReplay,...)

5

## PCAP Files

- PCAP – **P**acket **C**apture
- libpcap file format
- Commonly used for packet capture/generation
- Format:

Global Header	Packet Header	Packet Data	Packet Header	Packet Data	Packet Header	Packet Data
---------------	---------------	-------------	---------------	-------------	---------------	-------------

- Global header: magic number, version, timezone, max length of packet, L2 type, etc.
- PCAP Packet header:

ts_sec	ts_usec	incl_len	orig_len
--------	---------	----------	----------

6

## PCAP Files – a one slide outline

- PCAP – **P**acket **C**APture
- libpcap file format
- Commonly used for packet capture/generation
- Format:

Global Header	Packet Header	Packet Data	Packet Header	Packet Data	Packet Header	Packet Data
---------------	---------------	-------------	---------------	-------------	---------------	-------------

- Global header: magic number, version, timezone, max length of packet, L2 type, etc.
- PCAP Packet header:

ts_sec	ts_usec	incl_len	orig_len
--------	---------	----------	----------

Networking and Systems Measurements (L50)

7

## TCP Replay

- Free, software-based
- Replays network traffic stored in pcap files
  - Not just TCP
  - (not just pcap)
- Included in Linux
- Packets are sent according to pcap file timestamps

Networking and Systems Measurements (L50)

8

## Software based traffic generators

- Traditional tools (e.g., D-ITG, trafgen):
  - Rely on the interface provided by the kernel for packet IO
- Modern tools (e.g., MoonGen, pktgen, zsend):
  - Use special frameworks which bypass the network stack of an OS
  - Optimized for high speed and low latency
  - Cost: compatibility and support for high-level features

Networking and Systems Measurements (L50)

9

## MoonGen (Lab 3)

- A packet generator for  $\geq 10$  Gbit/s Ethernet
- Uses DPDK (an Intel originated idea – optimised for CPU consumption)
  - A set of libraries and drivers for fast packet processing
- (Sub-)microsecond timestamp accuracy
  - Using the NIC
- Rate control

Networking and Systems Measurements (L50)

10

## OSNT (Labs 2+3)

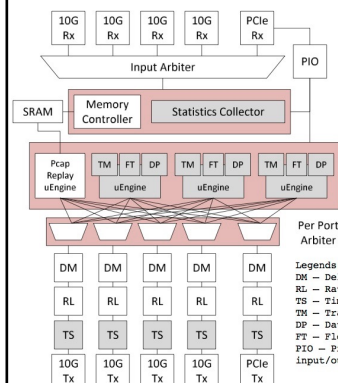


- Open source hardware/software traffic generator and capture system
- Built on top of NetFPGA platform
- Traffic generation using pcap file (currently)
- Rate controlled in hardware
- ~6ns resolution
- Recall (for 10Gb/s)
  - 10bits = 1ns and 1us  $\cong$  1kByte packet

Networking and Systems Measurements (L50)

11

## OSNT-TG architecture



- DRAM/SRAM used to store the packets
- DM and RL guarantee the output packet rate is the one assigned by the user

Legends:  
 DM – Delay Module  
 RL – Rate Limiter  
 TS – Time Stamp  
 TM – Traffic Model  
 DP – Data Pattern  
 FT – Flow Table  
 PIO – Programmed input/output

12

## High End Tools

- Cost from 1K's to 100K's of \$
- Typically hardware based
- With many software packages
- Scale to 400Gbit/second (per port)
- Accuracy: <1ns and many, (many,) *features*(\*)

(\* *features* (extra \$\$\$\$) may include: human readable output, 1982 era traffic model, trivial network model, no tracing, data samples smaller than 100 entries, first order statistics only, packet capture – any packet capture, synchronized clocks, ....)

13

## How to capture traffic?

- When did the packet arrive?
  - A hard question!
- Can part / all of the packet be captured?
- How many packets can be captured?
- What is the maximal rate of packets that can be captured?
- ...

14


## What is the time?

- Free running clocks, e.g.,
  - CPU's time stamp counter (TSC)
  - NIC's on board oscillator
  - Clocks drift!
- Synchronization signals, e.g.,
  - 1 PPS (pulse-per-second)
- Synchronization protocols, e.g.,
  - Network Time Protocol (NTP) – milliseconds accuracy
  - Precision Time Protocol (PTP) – microseconds accuracy (nanoseconds, depending on deployment)

15

### Host

## Timestamping



- At the port – highest accuracy
  - If you want to measure *the network*
- At the NIC – less accurate
  - Buffering, clock domain crossing etc.
- At the OS
  - Exhibits PCIe effects, scheduling dependencies
- At the Application – least accurate
  - Unless you are interested in the user's perspective – then it's the **only** place

16

## Traffic Capture



\$\$\$\$\$, Hardware, high quality  
(Ixia, Spirent,...)

\$\$ Software/hardware based, medium quality  
(DAG, OSNT, NIC based,...)

Commodity, Software, low quality  
(tcpdump, tshark, wireshark,...)

17

## tcpdump (libpcap)

- Software only
- libpcap (historically tcpdump)
- Other applications: tshark, wireshark...
- Captures data and <does stuff> including write stuff to a file
- Uses the pcap format (and others...)
- Timestamp comes from the Linux network stack (default: kernel clock)

18

## Packet Capture


Common example:

- `$ sudo tcpdump -i en0 -tt -nn host www.cl.cam.ac.uk`
- ```
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on en0, link-type EN10MB (Ethernet), capture size 65535 bytes
1507838714.207271 IP 192.168.1.107.50650 > 128.232.0.20.80: Flags [S], seq 3761395339, win 65535, options [mas 1460,nop,wscale 5,nop,nop,TS val 256908862 ecr 0,sackOK,eol], length 0
1507838714.207736 IP 192.168.1.107.50651 > 128.232.0.20.80: Flags [S], seq 527865303, win 65535, options [mas 1460,nop,wscale 5,nop,nop,TS val 256908862 ecr 0,sackOK,eol], length 0
...
```

19

## Where do I trace?

- Sometimes on the interface of a host (eg 'eth0')
  - `Tcpdump -i en1` # this will spew entries to the console one line per packet approximately
  - `-tt -nn` # useful options long form timestamps & numbers not names
- Interception using "Tap" (think wire-tapping)




20

## Endace (DAG)

- DAG - Data Acquisition and Generation
- A commercial data capture card
- Packet capture at line rate
- Timestamping in the hardware (at the port)
- Nanosecond resolution
- Clock synchronization possible
- Will be used in the labs

21




```
erf.binary dec
....0001 232ps,
....0010 466ps,
....0011 698ps,
....0100 931ps,
....0101 1163ps,
....0110 1397ps,
....0111 1629ps,
....1000 1862ps
erf = extensible record format
```

Why 232ps?

22

## NTP

- Designed for Internet-scale synchronization
  - E.g., email sent time < email received time
  - Milliseconds scale – emphasises frequency not phase
- A hierarchical system
- Using a few reference clocks
- Typically:
  - Host polls a few servers
  - Compensates for RTT and time offset
  - NTPv4 – RFC5905




23

## PTP

- IEEE standard 1588 (v2 – 1588-2008)
- Designed for local systems
  - Microsecond level accuracy or better
- Uses a hierarchical master-slave architecture for clock distribution
  - Grandmaster – root timing reference clock
  - Boundary clock – has multiple network connections, can synchronize different segments
  - Ordinary clock – has a single network connection (can be master or slave)
- (And many more details)

24



25

### Using NIC

- Either implement PTP-derived timestamps or just timestamp the packets sometimes in hardware most times... not...  
Not all NICs support time stamping
- Result: captured packets include a timestamp
- If PTP is used, end hosts are synchronized
- Else – free running counter

Networking and Systems Measurements (L50) 26

26

### Capturing to disk.....

- Most (physical) disk systems can not capture many 10's of Gb/s of data
- Capture takes resources!
- Format wars.... PCAP vs PCAP-ng vs others
- Binary representations / digital representations

Networking and Systems Measurements (L50) 27

27

### What makes high-speed capture hard?

- Disk bandwidth
- Host bandwidth (memory, CPU, PCIe)
- Data management
- Lousy OS and software APIs
  - Byte primitives are dreadful when you want information on events, packets, & transactions...
  - A lot of effort has been invested into reinventing ring-buffers (circular buffers) to accelerate network interface cards.
  - Performance networking was done for capture first....

Networking and Systems Measurements (L50) 28

28

### What makes high-speed capture work (better)?

- NVMe Disks
- Big machines, latest interfaces
- Collect metadata (version OS/system/hw/DNS)
- Bypass the OS
  - Older dedicated capture cards (e.g., Endace) pioneered kernel bypass capture
  - Any modern NIC 10Gb/s uses tricks that are useful for capture too

Networking and Systems Measurements (L50) 29

29

### Measuring Latency – Do's and Don't

- Make sure that you capture correctly
  - Disk, PCIe/DMA and other bottlenecks
- Make sure that your measurement does not affect the results
  - E.g., separate the capture unit from the device under test
- Understand what you are measuring
  - E.g., single host, application-to-application, network device etc.
- Make sure your traffic generator does not affect the results

Networking and Systems Measurements (L50) 30

30

## perf (not to be confused with iperf)

- So far we discussed *performance*
- What about *events*?
- Perf is a Linux profiler tool
- Allows us to instrument CPU performance counters, tracepoints and probes (kernel, user)

31

## perf

- list – find events
- stat – count events
- record – write event data to a file
- report – browse summary
- script – event dump for post processing

32

## Perf - example

```
i-/.ssh$ perf stat ps
PID TTY          TIME CMD
8747 pts/2        00:00:00 bash
11667 pts/2      00:00:00 perf
11670 pts/2      00:00:00 ps

Performance counter stats for 'ps':

   12.745507 task-clock (msec)    # 0.929 CPUs utilized
         4 context-switches      # 0.314 K/sec
         0 cpu-migrations        # 0.000 K/sec
       140 page-faults          # 0.011 M/sec
 32,322,489 cycles                # 2.536 GHz                    (40.80%)
<not supported> stalled-cycles-frontend
<not supported> stalled-cycles-backend
 27,644,922 instructions         # 0.86 insns per cycle        (68.86%)
  5,133,583 branches             # 402.776 M/sec              (68.92%)
   157,503 branch-misses        # 3.07% of all branches      (94.06%)

0.013726555 seconds time elapsed
```


↑  
the tool **scales** the count based on  
total time enabled vs time running

33

## Flame Graphs: an example of very clever visualization

- Parsing traces is like finding a needle in a haystack
- Flame graphs - Visualise the outputs of profiling tools
  - E.g., using perf, dtrace
- Easy to understand
- Open source
  - <https://github.com/brendangregg/FlameGraph>
  - Brendan Gregg has several other useful performance-related tools


34



35

## Flame Graphs

- Width is relative to “how much running on the CPU”
- Top-down shows ancestry
- Not good for idles – so don’t try to use for profiling network events!
- Different types of flame graphs
  - E.g. CPU, memory, differential



36

## Conclusion

- There are many .... So so many .... tools  
each is shaped by its heritage
- Select carefully (understand the limitations)
- Consider and collect metadata – always
- How will you find/process/interpret/visualize your data?